

Robotics for AI

Development lecture 3: Object Recognition

Table of content


- Deep Learning with Keras/Tensorflow
- Data gathering
- Data augmentation
- Numpy
- Behaviours
- Platform Instructions

Deep Learning

- 2D Convolutions → feature extraction
- Pooling layers → downscaling feature maps (hence less memory needed)
- Multilayer perceptron → classification


Deep Learning

- 2D Convolutions:
 - Kernel shape
 - Stride
 - Padding
 - Activation functions
 - sigmoid
 - relu
 - Tanh
 - More exist..



Deep Learning

- Pooling layers:
 - Max pooling
 - Average pooling
 - Other option for downsizing:
 - strides > 1 during convolutions



Deep Learning

We have a 30 x 30 x 32 featuremap.

If we want to go to classification / towards an MLP, we have to flatten:

$[30 \times 30 \times 32]$ featuremap $\rightarrow [30 * 30 * 32] = [28800]$

This step can be a huge memory bottleneck!

Deep Learning

So now we have [28800] feature vector.

If the next fully connected layer has 200 units, the required weight matrix will have the shape: [200 x 28800], resulting in 5 760 000 trainable weights!

Make sure your feature maps aren't too large before flattening.

Deep Learning with Tensorflow and Numpy

So, feature maps and images have a shape of [width, height, channels]

Having multiple images, you will have to use this shape:

[nImages, width, height, channels]

Hence, having 100 images with the size 1024 x 768 and 3 channels (r,g,b):

[100 x 1024 x 768 x 3]

Having 100 output vectors from the CNN, will result in a shape:

[100 x nClasses]

Deep Learning with Tensorflow and Numpy

You can train / run your CNN on batches of data.

Hence:

Input: [batch size, width, height, channels]

Output: [batch size, nClasses]


The model will be trained using the average error.

Data gathering:

In simulation:

- Spawn object (A script is provided to do this)
- Drive to and approach the objects
- Repeat:
 - Collect the ROI (we provide a tool to do this)
 - Respawn an object

For real objects, it's already done for you



Data Augmentation



Croppings are not perfect!

For every ROI, make slightly other ROIs from this object.

- Add a bit to a side / remove a bit from a side
- Slightly larger / slightly smaller ROIs

Lighting conditions are never constant in real life!

- Convert the image to HSV (hue, saturation and value) color space
- Use slightly different Value / Saturation values to augment brightness
- Convert back to RGB color space



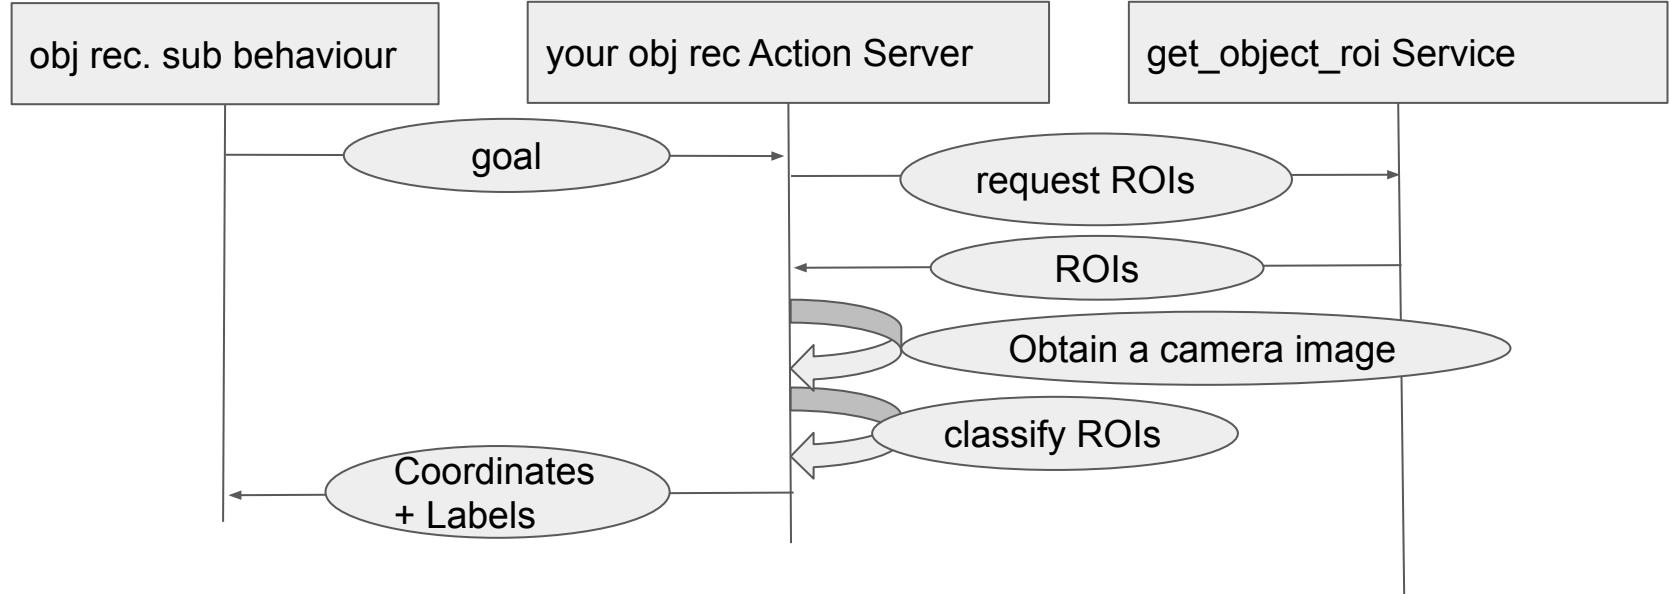
Numpy tutorial

Switching to other file to present...

Deep Learning with Tensorflow:


Tutorial: <https://www.tensorflow.org/tutorials/images/cnn>

Behaviour integration



Platform Instructions


- Turning on the platform
- Running programs on the platform.
- File transfer
- Turning off the platform



Platform Instructions

Power buttons:


- Main power switch
- Jetson
- USB
- Arm
- Motors



Platform Instructions

Power buttons:

- First set all buttons to *off*.
- Enable the main power switch.
- Up to object recognition:
 - Turn on the Jetson
 - Turn on the USB
 - Turn on the Motors



Platform Instructions

SSH usage:

- `ssh-copy-id roboticsX@nano`
 - This will make sure that you don't have to type your password at every SSH login.
- `ssh roboticsX@nano`
 - This gives you a terminal on the robot

Platform Instructions

Cloning and building

- Clone your repository in `catkin_ws/src`
- Build your repository with the following command in `catkin_ws/src`:
 - `catkin build -j2`

Platform Instructions

ROS

- Run *roscore* on the robots
- To allow a pc terminal to communicate with the robot, use the following command in every terminal that should communicate with the robot, e.g. to use RViz:
 - **export ROS_MASTER_URI=http://nano:11311**

Platform Instructions

Start up

To start all sensors and controllers, run ***on the robot:***

roslaunch robot bringup_robot.launch

To start the arm controller, run ***on the robot:***

roslaunch dynamixel_arm_controller arm_controller.launch

To start navigation, assuming you have the map, run ***on the robot:***

roslaunch navigation real_world.launch

Platform Instructions

To transfer your CNN model to the robot:

```
scp -r folder_name roboticsX@nano:~/data/
```

Platform Instructions

To shutdown the robot:

sudo poweroff

Wait until the green LED of the Jetson is off, then power-off the robot



That's all Folks!